

19

CONSCIOUSNESS AND INTENTIONALITY

David Pitt

A mental state is conscious just in case there is something it is like to be in it. The properties in virtue of which there is something it is like to be in a mental state are *phenomenal properties*, or *qualia*. A mental state is intentional just in case it is *about* something, and thereby has truth or veridicality conditions. The feature of an intentional state in virtue of which it has these properties is called its *intentional content*. In analytic philosophy of mind there was for many years a consensus that consciousness and intentionality are properties of metaphysically exclusive kinds. Conscious qualitative states, such as visual, auditory and olfactory experiences, do not, *per se*, have intentional content; and intentional states, such as thoughts, beliefs, desires and intentions do not, *qua intentional*, have phenomenal properties. To be sure, perceptual states such as seeing a dog or hearing it bark are perceptions *of* dogs and barks, and thereby have intentional content. But their intentionality was typically taken to be determined by causal relations between perceiver and perceived, and not by any intrinsic qualitative features they might have. And though thoughts, beliefs and desire may be conscious, whatever qualitative features might be associated with thinking, believing and desiring were taken to be irrelevant to their intentional content. In general, the phenomenal character of conscious states was seen as having no essential connection to their intentional contents.

Consciousness is extremely difficult (some think impossible) to explain within the naturalist framework that has prevailed in analytic philosophy of mind for most of the twentieth century, and into the twenty-first. Intentionality, on the other hand, insofar as it is a phenomenon that is not essentially tied to consciousness, was seen to be more tractable, and various theories grounding it in or reducing it to natural relations between the brain and the world it represents were proposed and developed. Philosophers working on intentionality, both perceptual and cognitive, felt they could safely ignore the vexing problem of the naturalization of consciousness.

More recently, however, this consensus has begun to weaken, as naturalistic theories of intentionality have faced problems that a growing number of philosophers believe are due to their failure to take conscious qualitative experience into account. These philosophers have argued that intentionality is essentially an experiential phenomenon, and, as such, cannot be reductively explained unless consciousness can – however problematic this may be for the naturalistic program in philosophy of mind. They have taken a stance reminiscent of classical phenomenology, which “brackets” the relation of experience to the world in order to study it on its own terms. These *analytic* phenomenologists tend to bracket the relation between experience and the brain,

pursuing a phenomenal theory of intentionality free from, as Charles Siewert (2011: 242) so memorably put it, “the tyrannizing anxieties and ambitions of mind-body metaphysics.” While not ignoring the metaphysical problem of consciousness, these analytic phenomenologists insist that reductive explanation is not the only project one might profitably pursue in the study of consciousness.

1 Causal-Informational Psychosemantics

Fred Dretske was set to be the Darwin of intentionality. His insight that causal relations, insofar as they carry *information* about the occurrence of the events they relate, establish a kind of proto-intentionality, is profound. It is the kind of idea – intuitive, simple and powerful – we all wish we had thought of (and wonder why we didn’t).¹ Though not yet what *we* have, this proto-intentionality is sufficiently *like* it to get us a conceptual foot in the seemingly unopenable door between this aspect of mind and our physical constitution. Dretske’s idea promised to show how it is *possible* that a highly sophisticated and puzzling aspect of our mental nature could arise from simple beginnings, by entirely natural processes.

In the 1980s and ’90s there was, understandably, a great deal of excitement among analytic philosophers of mind over this idea. Jerry Fodor went as far as to suggest that (modulo a syntactic solution to Frege’s Puzzle) “Turing and Dretske have between them solved the mind/body problem” (Fodor 1994: 56). Turing showed how a physical thing could *reason*, and Dretske showed how a physical thing could *represent*. The philosophical project of naturalizing the mind, of bringing it within the scope of the kind of empirical methodology that led to such spectacular successes in our understanding of the world, seemed to be, if not complete, at least halfway there.

The view has the added benefit of building a connection between thought and its objects into the very nature of representational content. Concepts are individuated by the object(s) or property instantiation(s) whose presence is lawfully causally correlated with their occurrence, and thus acquire their contents and their extensions simultaneously.

There was (as Dretske and Fodor were always well aware) still the problem of consciousness to be addressed. Causal relations *per se* do not seem to be sufficient to bring about conscious experience, or even some kind of proto-conscious experience. Qualia freaks would have to await their own Darwin. But the other half of the mind-body problem was, from a philosophical point of view, in its essential outlines, thought to have been solved.

Philosophy being philosophy, of course there were dissenters all along. In particular, there have been those, such as John Searle and Galen Strawson, who have long insisted that *genuine* intentionality (what *we* have) is essentially a conscious, experiential phenomenon. Searle has argued for what he calls the “connection principle” (Searle 1992), according to which a mental state cannot have fine-grained intentional content (what he calls “aspectual shape”) unless it is either conscious or potentially conscious,² and Strawson (1994) has argued for the essential experientiality of mentality in general, and of conceptual intentionality in particular. According to these theorists, resources sufficient for constructing an information processor are not sufficient for constructing a mind, since information *per se* is not conscious, and consciousness is required for genuine intentionality. Another important defender of this idea is Charles Siewert (1998).

Causal-informational theorists have, unsurprisingly, resisted this claim. If true, it would short-circuit their naturalistic explanation of intentionality, since at present there is no adequate naturalistic account of conscious experience (and some would argue that there can never be one). Fodor even pronounced commitment to an essential link between intentionality and conscious experience “intellectual suicide.”³ But, as we will see, it is a position that has recently

been gaining adherents in analytic philosophy of mind, who so far appear to have remained intellectually above ground.

In spite of their promise, causal-informational theories face internal difficulties – the most persistent of which have been problems of indeterminacy. There is Quine's Problem, which arises out of what may be called *causal superimposition*; the Disjunction Problem, which arises out of what may be called *causal spread*; and the Stopping Problem, which arises out of what may be called *causal depth*. In all of these cases, there are multiple candidates for content determiner/extension, and no obvious way to choose among them derivable from the basic machinery of the theory.

Quinean examples of indeterminacy of radical translation (Quine 1960) can be taken to show that for any property that is a candidate for determining the content of a concept (the meaning of a term), there are indefinitely many other simultaneously instantiated (superimposed) properties that cannot be teased apart causally. Any instantiation of *rabbithood*, for example, is also, necessarily, an instantiation of *undetached-rabbit-parts-hood*, *rabbit-stage-hood*, and indefinitely many other properties. Assuming that these properties are distinct, they are candidates for distinct contents for the meaning of 'rabbit' (and the concept [mental representation] RABBIT). (Names for concepts are here written in SMALL CAPS, and names of properties in *italics*.) Given that these properties are (at least physically) necessarily instantiated by the same things, there can be no lawful relations between mental states and *one* of them that are not also lawful relations between mental states and *all* of them. Hence, a causal-informational theory cannot, at least *prima facie*, assign one of them as the content of RABBIT. There is by the theory's lights no fact of the matter about which of these properties is content-determinative of the concept RABBIT (the term 'rabbit').

Though Quinean examples can be taken to show as entailing *indeterminacy* of content, they can also be viewed as entailing massive *disjunctiveness* of content. On this construal, the content of RABBIT would be *rabbithood or undetached-rabbit-parts-hood or rabbit-stage-hood or* In this case there would be a fact of the matter about what the content of a given concept is, but it would be, counterintuitively, open-endedly disjunctive. This is problematic because, as Fodor has often pointed out (e.g., Fodor 1987), there ought to be psychological generalizations that apply to mental states in virtue of their content. However, in keeping with the naturalistic project, such laws would be *causal* (or otherwise nomological). But natural laws typically are not formulated in terms of disjunctive properties, which do not in general constitute *natural kinds*.

Dretske (1981) himself recognized this problem (named the "Disjunction Problem" in Fodor 1984), which arises from the fact that there are causal correlations between the occurrence of mental representations and the presence of a wide range of things (property instantiations) that are, intuitively, not in the extension of those representations. Thus, though there may be a law-like regularity between horses (instantiations of *horsehood*) and occurrences of the concept HORSE, such relations *also* hold between HORSE occurrences and indefinitely many other things: *donkeys on dark nights, zebras in the mist, merest ripples in horse-infested waters*,⁴ ... – anything that might cause one to think, correctly or incorrectly, (e.g.) LO, A HORSE! Thus, for HORSE (or any empirical concept), there is a *spread* of different property instantiations (by distinct objects) sufficient for its tokening, and, hence, by the theory's lights, sufficient for determining its content. But HORSE cannot mean *all* of these indefinitely many things. And the reasons for resisting a *disjunctive* content are the same here as they were in the causal superimposition cases.

Indeed, though this is not always remarked upon, one could just as well construe this as a problem of indeterminacy: there is, consistent with the resources of the theory, no fact of the matter about which one of the indefinitely many causally correlated property instantiations determine a concept's content.

Another problem (named the "Stopping Problem" in Strawson 2008) first arises when the causal relations that are supposed to establish content hold between mental states and *distal*

objects.⁵ Thus, causal relations to cows – instantiations of *cowhood* – are supposed to constitute a mental representation the concept COW. But there are *also* causal relations between occurrences of COW and any link in the causal chain between cows and COWs. These include links *within* the perceptual system, such as *bovine retinal images*, *bovine olfactory bulb stimulations*, *bovine visual or olfactory cortex activation patterns*, etc.,⁶ as well as links between retinal images (or other sensory-organ representations) and cows – such as *cow reflections*, *cow shadows*, *cow breezes*, There are also less obvious candidates, like *photons reflected from a cow*, *the cow's parents*, *distant ancestor bovine species*, ..., *the Big Bang*. All of these can lay equal claim to inclusion in the causal chain leading to tokenings of COW, although, obviously, the vast majority of them are not plausible candidates for being (or determining) the content or extension of the concept COW.

The causal chains connecting concept tokenings to their content-conferring property instantiations are *deep*, involving a densely packed series of property instantiations (events) as links. And while we may find it impossible to take seriously candidates such as objects or events in the distant past, or property instantiations undetectable by us, if *all* we have at our disposal is causal relations, it is not obvious what principled reasons there could be for excluding any of them. And if there is no way to prune away the unwanted causes, then we are faced, as with the other problematic cases, with the invidious choice between indeterminacy and massive disjunction.⁷

And there are other apparent problems, as well: How are causal theories to explain the contents of mathematical, logical and other concepts, whose referents are abstract, causally-inert objects? Or the contents of concepts of non-existent objects?

Causal-informational theorists have expended considerable effort and ingenuity in the search for a solution to these problems (see e.g. Dretske 1988, 1995; Fodor 1987, 1990; Millikan 1984, 1989; Neander 1995; Papineau 1998; Prinz 2002; Rupert 1999, to cite just a few examples from a very large literature). Some see a solution in teleology – the evolved function of representation-producing mechanisms; though there are residual indeterminacy problems for such views (see Fodor 1990). Others appeal to causal-inferential relations among mental representations (see Block 1986; Field 1977; Harman 1973, 1987; Loar 1981; and McGinn 1982 for foundational statements of the view). These “conceptual-,” “functional-,” or “inferential-role” theories are typically integrated with Dretske-style accounts in constructing “two-factor” (internal and external, “narrow” and “wide”) theories of content. These theories have their own technical difficulties, arising from their *prima facie* commitment to meaning holism (see e.g. Fodor and Lepore (1992). (An intuitive objection to such views is that inferential relations among concepts are determined by their contents, not vice versa.) But it would not be accurate to say that naturalistic approaches of these kinds are defunct.⁸

2 Phenomenal Intentionality

Other philosophers have proposed that in order to solve these problems – or, even better, to avoid them entirely – causal relations should be replaced with (or at the very least supplemented by) qualitative features of experience as determiners of content. Searle and Strawson have already been mentioned as early analytic proponents of an experience-based approach to intentionality.⁹ Searle (1987) responds to Quinean indeterminacy; and Strawson addresses the Stopping Problem in his 2008. It has also been argued that phenomenology can solve the Disjunction Problem (Pitt 2009; Horgan and Graham 2012).

The shared idea is that what our concepts are concepts *of* is what we *take them to be of*, where *taking* is a manner of *experiencing*. What HORSE means is what *we* mean by it; and what we mean is experiential, and introspectively available to us. We *know*, from a first-person perspective, that the extension of HORSE is horses, and not *horse-part-fusions* or *zebras in the mist* or *equine retinal*

arrays,¹⁰ And we know this in this way because conceptual contents (and thought contents) are *experiential* in nature.

Searle calls the experiential content of a concept its “aspectual shape.” Strawson (1994) calls it “understanding experience.” Siewert (2011) speaks of “phenomenal thought.”¹¹ It has lately come to be known as “cognitive phenomenology” (Pitt 2004; Bayne and Montague 2011a; Chudnoff 2015a; see Strawson 1986 for an early use of this term). Without claiming that everyone who subscribes to this view agrees about the nature of conceptual experience and its relation to intentional mental content (some theorists claim that it does not determine content at all [Siewert 1998], some say it constitutes only an internally determined component of content [Horgan and Kriegel 2008; Strawson 2008], while others reject the idea that content should be factored into internally and externally determined components [Pitt 2013]), we can say that there is a shared commitment to the thesis that genuine conceptual intentionality of the kind we have is essentially an experiential phenomenon. Without experience (which for most philosophers means without *consciousness*) there can be no mental representation with the fineness of grain or selectivity that our thoughts and concepts display.

Apart from its value as a prophylactic (or cure) for Indeterministic Disjunctivitis, conceptual phenomenology has been recommended on independent grounds.

One common form of argument is from *phenomenal contrast*. In one kind of case, we are invited to compare the experience of hearing discourse in a language that is understood to the experience of discourse in a language that is not understood (Strawson 1994: 5–6). In another, we are invited to consider *changes* in our own conscious occurrent thought (Siewert 1998: 275–278). In yet another, we are to imagine an individual who lacks all sensory, emotional, algedonic, etc., experience, yet who can still *think*, and consider what it is like for this individual to reason mathematically (Kriegel 2015: 56–62). In all cases, it is argued that there is a phenomenal difference, a difference in what it’s like for the thinker, and, further, that this is not a difference in familiar kinds of phenomenology, such as that of verbal or auditory imagery, emotional tone, etc. It is then concluded that there is an irreducible, distinctively *cognitive* kind of experience that accompanies (or constitutes) thinking, differences in which account for the experiential contrasts.¹²

Phenomenal contrast arguments are vulnerable to competing claims about what the contrast between experiences with and without understanding actually consists in. What proponents attribute to a difference in cognitive phenomenology, critics maintain is a difference in auditory, visual, emotional, or some other more familiar kind of phenomenology. Such positions are bolstered by claims of a lack of introspective evidence in the objector’s own experience for the existence of such *sui generis* cognitive phenomenology.¹³ Disputes over what is phenomenally manifest in introspection are notoriously difficult (though not impossible) to adjudicate. This has led some to doubt whether the phenomenal contrast strategy is the best way to try to establish the existence of cognitive phenomenology. (Sacchi and Voltolini [2016] offer a version of the contrast argument that, they claim, does not rely on introspection.)

A different sort of approach, due to Strawson, focuses on the *significance* or *value* of conscious experience in general, and of conscious thought in particular. Strawson (2011) argues that our conscious experience would be significantly *less interesting* if it did not include an experience of thinking. If thoughts were just unconscious subpersonal computational states, our conscious mental lives would be drastically impoverished. We would have no experience of *grasping* truths, of *wondering* why something is the case, of *realizing* and *solving* problems, etc.

Another type of argument for cognitive phenomenology appeals to a particular kind of self-knowledge we are capable of. Pitt (2004) argues that it is possible to know, consciously, introspectively and non-inferentially, what one is consciously occurrently thinking, and that this would not be possible if thought (and conceptual) contents were not immediately present in

consciousness. Just as one can know in this way that one is hungry, hearing a trumpet or tasting ashes, because there is something it is like to be in these states, one can know *that* one is thinking, and *what* one is thinking, because there is something it is like to think, and what it is like to think thoughts with different contents is phenomenally different. Conscious occurrent thoughts could not be introspectively distinguished from other kinds of conscious states, and from each other, in this way if they were not phenomenally individuated. Moreover, since it is possible to have auditory or visual experience of linguistic expressions without thinking what they mean, or thinking anything at all, this individuating phenomenology cannot be the phenomenology of inner speech or visualization.

Pitt (2009) argues, further, that this cognitive kind of phenomenology *is* cognitive intentional content. To consciously think that three is a prime number is to consciously token a maximally determinate cognitive-phenomenal type which *is* the proposition that three is a prime number. (Just as to be in a maximally determinate pain state is to token a maximally determinate pain type.)

Pitt (2011) offers another argument for cognitive phenomenology, based upon the claim that conscious states, as such, are individuated phenomenologically. That is, what distinguishes conscious states of different kinds is their different kinds of phenomenal character. Conscious sensory states, such as visual, auditory and olfactory experiences, are distinguished by, respectively, visual, auditory and olfactory phenomenology, each a *sui generis* kind of experiential quality. And conscious sensory states within a given modality are, as such, individuated by different determinate phenomenologies within their respective determinable phenomenal kinds. Pitt argues that conscious thought, qua conscious, is individuated in the same way as other kinds of conscious experience, as are distinct thoughts within the cognitive experiential modality. Hence, there must be a proprietary, distinctive and individuating phenomenology of occurrent conscious thought.

Perceptual states are also intentional. In their various modalities, they represent to us the world around us, providing information about the existence and states of the things with which we interact. And they can be more or less accurate, veridical or not. What is the role of consciousness in the intentionality of perception? Obviously, conscious perceptual experiences must be conscious. But what role do the phenomenal properties apparent in conscious experience play in determining the intentional content of a perceptual state – what it is a perception *of*? On what can be called the Pure Causal View, they play no role whatever. A perceptual state is a representation of an object or property (instantiation) if and only if it is caused by that object or property. Whatever qualitative properties may be consciously apparent determine, at best, only how accurately or inaccurately a perceptual state represents, not whether or not it represents. Toward the other end is what Montague (2016) calls the Matching View, according to which there is a (probably vague) limit to how badly a perceptual state can misrepresent its cause before it ceases to be a perception of it.

Most (if not all) philosophers would agree that a causal relation between token perceptual states and specific objects or properties is *necessary* for genuine perception. No state not caused by an elephant is a perception *of* an elephant. The role of causation with respect to perceptual states is thus different from its role with respect to cognitive (conceptual) states. In the latter case, we want to allow that token concepts can be *of* things that are *not* their token causes. A token concept ELEPHANT should be a concept *of* elephants (have elephants in its extension), no matter what causes it, and whether or not it was caused by any external thing or property. But a token perceptual state cannot be a perception of an elephant unless it is caused by an elephant. Because of this difference, the Disjunction Problem does not arise for perceptual states. Perceptions of elephants cannot be caused by hippos-in-the-mist or large grey rocks, or by nothing at all.

Quine's problem also does not arise for perceptual states, since, for example, a perceptual state caused by an elephant is also caused by an elephant-stage and a sum of undetached elephant parts, etc. The *conceptual* distinctions do not seem to be relevant to what is being perceived in the way they are relevant to what is being thought about.

But the Stopping Problem does arise. Any state caused by an *F* is also caused by other links in the causal chain leading to the occurrence of the state. A visual perception of an elephant is caused by the elephant; but it is also caused by whatever caused the elephant, the photons reflected from the elephant, the firing of cells in the retina, the lateral geniculate nuclei and the primary visual cortex, etc. – none of which we would want to say the experience is *of*.

The Matching View has a straightforward solution to this problem: the visual experience one has upon looking at an elephant is not an experience of any of these other causes because it does not *resemble* any of them. This is analogous to the cognitive-phenomenal solution to the Quine and Disjunction Problems for conceptual representations – the concepts RABBIT and RABBIT-STAGE, HORSE and COW-IN-THE-MIST, are introspectively distinguishable cognitive experiences. What it is like to think that something is a rabbit is different from what it is like to think that it is a rabbit-stage.

Some philosophers (e.g. Evans 1982 and Dretske 2006) have argued that in order for a state to be a perception of an *F*, it must not just be caused by an *F*, but also enable the perceiver to *locate* or *track* the *F*. And this might seem to be enough to solve the perceptual Stopping Problem, since the state of perceiving an elephant does not provide the perceiver information about the location of the elephant's ancestors, the photons bouncing off it, the perceiver's retina, or parts of the perceiver's brain. Moreover, since on this account the state itself need not (at least for Dretske) be phenomenally conscious, it need not *resemble* its cause in any robust sense. And even if it is acknowledged that perceptual states are typically conscious, and that conscious states have (or present) qualitative properties, one may allow that these properties establish whether or not the state resembles its cause, but still deny that resemblance is necessary for genuine perception.

Montague insists, however, that there are limits to how badly a conscious perceptual state can misrepresent its cause before it is disqualified as a perception of it. On her Matching View, a perceptual state “must represent a sufficient number of [an] object's properties correctly in order for it to be true that one [perceives] it” (Montague 2016: 156). On this view, an experience that in no way resembles an elephant cannot be a perception of the elephant that caused it.

The intuitions on both sides are respectable. On the one hand, it seems reasonable to say that an experience caused by an *F* is a perception of that *F* no matter how unlike its cause it is – just as it seems reasonable to say that a photograph is of an *F* if it was photons bouncing off the *F* that were responsible for its production (cf. Evans 1982, 78), no matter how little it resembles its cause; or a painting is a painting of an *F* if the artist intended it to be a painting of an *F*, no matter how little it might resemble the *F* (cf. modern art).

On the other hand, if we consider things from the perspective of the representation itself, it seems reasonable to say that resemblance *is* required. No one shown a picture of an elephant would take it to be a picture of a pink Cadillac, or vice versa. And no one would take a completely blank image to be a photograph of either an elephant or a pink Cadillac. Moreover, it seems entirely natural to say that an image with the appropriate properties is an image of an elephant, whether or not it resulted from causal interaction with one, and somewhat perverse to say that such an image is not an image of an elephant, because it was not caused by one.

These intuitions are not inconsistent. There is a perfectly good sense of ‘a perception *of* an *F*’ on which it means *a perception caused by an F*, and an equally good sense on which it means *a perception resembling an F*. The latter sense is commonly marked out with the phrase ‘perception *as*

of an *F*' (or of an *F* as an *F*). A perception of an *F* (like a photograph or picture of an *F*) may or may not also be a perception *as of* an *F*. Being caused by an *F* does not entail, and is not entailed by, resembling an *F*. A state caused by an elephant could resemble virtually anything, or nothing at all; and a state resembling an elephant could be caused by virtually anything, or nothing at all. (Additionally, the former sense may be used in reference to a perception [or photograph or painting] of a *particular F*, the latter in reference to a perception [or photograph or painting] of a *typical F*, though none in particular.)

However, if the issue is the intentionality of perceptual experience itself, then it is arguable that the latter sense of 'perception of' is more appropriate. For the content of perceptual experience *as one has it* is constituted by its phenomenal character. Perceivers do not have direct access to external causal relations between objects and their perceptions of them. And if the *role* of perception is to inform perceivers of the existence and states of external objects, then complete misrepresentation of its external cause should disqualify an experience as a genuine perception, since such an experience would be (more or less) *useless* to the perceiver.

Dretske and others (e.g. Dretske 1995; Harman 1990; Lycan 1996; Tye 2000) have proposed extensions of the causal-informational theory to give a naturalistic account of the qualitative properties apparent to us in perceptual experience. Such "reductive representationalist" views (see Chalmers 2004 for terminological clarification) attempt to explain the phenomenology of perception in terms of causal-informational representation of objectively instantiated phenomenal properties. The yellowness one might mention in describing what it is like to see a ripe banana, for instance, is a property of the banana, not one's experience of it. And it is easy to see how this account could be used to solve the Stopping Problem for perception: a perceptual state represents the thing whose phenomenal properties are apparent to the perceiver. However, this "qualia externalism" (see Byrne and Tye 2006) faces serious problems in accounting for dreams, illusions and hallucinations (Thompson 2008; Pitt 2017). (Moreover, it is far from obvious how externalist theories of this kind could solve the indeterminacy problems for cognitive states. See Byrne 2011, 2008 and Pitt 2011.)

3 Conclusion

There is a common point to be made about the role of phenomenology in determining conceptual and perceptual intentionality (content). A theory that takes causal-informational relations between representation and represented to be sufficient to determine the content of the representation (what the representation is about) will encounter indeterminacy/disjunction problems that cannot be solved in purely causal-informational terms. The diagnosis offered by advocates of phenomenal intentionality is that such difficulties are bound to arise whenever the intrinsic properties of representations are ignored. Such properties have an essential role in both determining representational contents and making them available to the thinker or perceiver. If thought and perception are to establish useful and accurate representational connections between conscious thinker-perceivers and their worlds, it must be apparent to them *what* is being represented and *how* it is being represented, and how a thing is represented must sufficiently resemble (accurately characterize) what it represents. In consciousness, appearance is, necessarily, phenomenal. Nothing can appear to a thinker-perceiver without appearing in some *way* or other, and the ways of appearing are constituted by phenomenal properties. And nothing can be accurately and (hence) usefully conceived or perceived unless the way it appears to the thinker-perceiver is the way it *is*. In spite of the fact that consciousness and phenomenality stubbornly resist naturalistic explanation, no theory of intentionality can afford to ignore them.

Notes

- 1 See Dretske (1981, 1988, 1995). C.B. Martin had a different, but also inspired, idea when he noticed that the relation between dispositions and their manifestations can also be seen as a kind of proto-intentionality. Dispositional states are directed at, indicate, or point to, their manifestations. See Martin (2008).
- 2 See Fodor and Lepore (1994) and Gennaro 2012 (sec. 2.3.1) for critical discussion of Searle's connection principle. Searle (1984) also objected to the idea that Turing solved the naturalization problem for reasoning, arguing that rule-governed symbol-manipulation without understanding (for Searle, a form of experience) is not thinking.
- 3 I am not aware of this remark appearing in print. I have heard Fodor say it, and Strawson reports it in 2008.
- 4 Cf. Fodor (1990).
- 5 As Strawson notes, this is a common problem for causal theories generally (e.g., the causal theory of perception, to be discussed below).
- 6 As pointed out also in Sterelny (1990), Antony and Levine (1991), Adams and Aizawa (1997) and Fodor (1990).
- 7 The Stopping Problem has also been called the "horizontal disjunction problem." The three problems discussed here are really versions of a general problem that we might call the "Problem of Causal Proliferation."
- 8 Philosophical views are rarely, if ever, definitively defunct. What usually happens is that people get bored with their problems and move on to something else. Often enough, old views get resurrected once new ones become stale.
- 9 In the Phenomenological tradition, the experiential nature of intentionality is taken to be self-evident.
- 10 Indeed, as has often been pointed out, if we could make no such distinctions as those between the contents *rabbit* and *rabbit-stage*, indeterminacy and disjunction would not appear to us to be *problems* at all. Of course, Quine famously denied that, after all, there is a difference between *rabbit* and *rabbit-stage* for the radical translator. But, as Searle (1987) argued, this strains credibility (to say the least). It seems more plausible to see Quinean indeterminacy as a *reductio* of empiricist semantics.
- 11 See also the discussion of the experience of thinking in Siewert (1998, ch. 8).
- 12 Other proponents of this kind of argument include Horgan and Graham (2012), Horgan and Tienson (2002), Moore (1962), Peacocke (1998) and Siewert (1998, 2011).
- 13 See, e.g., Carruthers and Veillet (2011), Chudnoff (2015b), Koksvik (2015), Levine (2011), Pautz (2013), Prinz (2011) and Tye and Wright (2011).

References

- Adams, F. R. and Aizawa, K. (1997) "Fodor's Asymmetric Causal Dependency Theory and Proximal Projections," *Southern Journal of Philosophy* 35: 433–437.
- Antony, L. and Levine, J. (1991) "The Nomic and the Robust," in B. M. Loewer and G. Rey (eds.), *Meaning in Mind: Fodor and His Critics*, Oxford: Blackwell.
- Bayne, T. and Montague, M. (eds.) (2011) *Cognitive Phenomenology*, Oxford: Oxford University Press.
- Bayne, T. and Montague, M. (2011a) "Cognitive Phenomenology: An Introduction," in Bayne and Montague (2011).
- Block, N. (1986) "Advertisement for a Semantics for Psychology," *Midwest Studies in Philosophy* 10, Minneapolis, MN: University of Minnesota Press.
- Byrne, A. (2008) "Knowing That I Am Thinking," in A. E. Hatzimoysis (ed.), *Self-Knowledge*, Oxford: Oxford University Press.
- Byrne, A. (2011) "Transparency, Belief, Intention," *Aristotelian Society Supplementary Volume* 85: 201–221.
- Byrne, A. and Tye, M. (2006) "Qualia Ain't in the Head," *Noûs* 40: 241–255.
- Carruthers, P. and Veillet, B. (2011) "The Case Against Cognitive Phenomenology," in Bayne and Montague (2011).
- Chalmers, D. (2004) "The Representational Character of Experience," in B. Leiter (ed.), *The Future for Philosophy*, Oxford: Oxford University Press.
- Chudnoff, E. (2015a) "Phenomenal Contrast Arguments for Cognitive Phenomenology," *Philosophy and Phenomenological Research* 90: 82–104.
- Chudnoff, E. (2015b) *Cognitive Phenomenology*, London: Routledge.

- Dretske, F. (1981) *Knowledge and the Flow of Information*, Cambridge, MA: MIT Press.
- Dretske, F. (1988) *Explaining Behavior*, Cambridge, MA: MIT Press.
- Dretske, F. (1995) *Naturalizing the Mind*, Cambridge, MA: MIT Press.
- Dretske, F. (2006) "Perception Without Awareness," in T. S. Gendler and J. Hawthorne (eds.), *Perceptual Experience*, Oxford: Oxford University Press.
- Evans, G. (1982) *The Varieties of Reference*, Oxford: Oxford University Press.
- Field, H. (1977) "Logic, Meaning and Conceptual Role," *Journal of Philosophy* 69: 379–409.
- Fodor, J. A. (1984) "Semantics, Wisconsin Style," *Synthese* 59: 231–250.
- Fodor, J. A. (1987) *Psychosemantics*, Cambridge, MA: MIT Press.
- Fodor, J. A. (1990) "A Theory of Content," in *A Theory of Content and Other Essays*, Cambridge, MA: MIT Press.
- Fodor, J. A. (1994) *The Elm and the Expert*, Cambridge, MA: MIT Press.
- Fodor, J. A. and Lepore, E. (1992) *Holism: A Shopper's Guide*, Oxford: Blackwell
- Fodor, J. A. and Lepore, E. (1994) "What Is the Connection Principle?" *Philosophy and Phenomenological Research* 54: 837–845.
- Gennaro, R. (2012) *The Consciousness Paradox*, Cambridge, MA: MIT Press.
- Harman, G. (1973) *Thought*, Princeton, NJ: Princeton University Press.
- Harman, G. (1987) "(Nonsolipsistic) Conceptual Role Semantics," in E. Lepore (ed.), *New Directions in Semantics*, London: Academic Press.
- Harman, G. (1990) "The Intrinsic Quality of Experience," *Philosophical Perspectives* 4: 31–52.
- Horgan, T. and Graham, G. (2012) "Phenomenal Intentionality and Content Determinacy," in R. Schantz (ed.), *Prospects for Meaning*, Berlin: de Gruyter.
- Horgan, T. and Kriegel, U. (2008) "Phenomenal Intentionality Meets the Extended Mind," *The Monist* 91: 347–373.
- Horgan, T. and Tienson, J. (2002) "The Intentionality of Phenomenology and the Phenomenology of Intentionality," in D. Chalmers (ed.), *Philosophy of Mind: Classical and Contemporary Readings*, Oxford: Oxford University Press.
- Kriegel, U. (ed.) (2013) *Phenomenal Intentionality: New Essays*, Oxford: Oxford University Press.
- Kriegel, U. (2015) *The Varieties of Consciousness*, Oxford: Oxford University Press.
- Koksvik, O. (2015) "Phenomenal Contrast: A Critique," *American Philosophical Quarterly* 52: 321–334.
- Levine, J. (2011) "On the Phenomenology of Thought," in Bayne and Montague (2011).
- Loar, B. (1981) *Mind and Meaning*, Cambridge, MA: Cambridge University Press.
- Lycan, W. G. (1996) *Consciousness and Experience*, Cambridge, MA: MIT Press.
- McGinn, C. (1982) "The Structure of Content," in A. Woodfield (ed.), *Thought and Object*, Oxford: Oxford University Press.
- Martin, C. B. (2008) *The Mind in Nature*, Oxford: Clarendon Press.
- Millikan, R. (1984) *Language, Thought and Other Biological Categories*, Cambridge, MA: MIT Press.
- Millikan, R. (1989) "Biosemantics," *Journal of Philosophy* 86: 281–297.
- Montague, M. (2016) *The Given*, Oxford: Oxford University Press.
- Moore, G. E. (1962) "Propositions," in *Some Main Problems of Philosophy*, New York: Collier Books: 66–87.
- Neander, K. (1995) "Misrepresenting and Malfunctioning," *Philosophical Studies* 79: 109–114.
- Papineau, D. (1998) "Teleosemantics and Indeterminacy," *Australasian Journal of Philosophy* 76: 1–14.
- Pautz, A. (2013) "Does Phenomenology Ground Mental Content?" in Kriegel (2013).
- Peacocke, C. (1998) "Conscious Attitudes, Attention, and Self-Knowledge," in C. Wright, B. C. Smith, and C. Macdonald (eds.), *Knowing Our Own Minds*, Oxford: Oxford University Press.
- Pitt, D. (2004) "The Phenomenology of Cognition, or, What Is It Like to Think That P?," *Philosophy and Phenomenological Research* 69: 1–36.
- Pitt, D. (2009) "Intentional Psychologism," *Philosophical Studies* 146: 117–138.
- Pitt, D. (2011) "Introspection, Phenomenality and the Availability of Intentional Content," in Bayne and Montague (2011).
- Pitt, D. (2013) "Indexical Thought," in U. Kriegel (2013).
- Pitt, D. (2017) "The Paraphenomenal Hypothesis," *Analysis* 77: 735–741.
- Prinz, J. (2002) *Furnishing the Mind: Concepts and Their Perceptual Basis*, Cambridge, MA: MIT Press.
- Prinz, J. (2011) "The Sensory Basis of Cognitive Phenomenology," in Bayne and Montague (2011).
- Quine, W. V. O. (1960) *Word and Object*, Cambridge, MA: MIT Press.
- Rupert, R. (1999) "The Best Test Theory of Extension: First Principles," *Mind and Language* 14: 321–355.
- Sacchi, E. and Voltolini, A. (2016) "Another Argument for Cognitive Phenomenology," *Rivista Internazionale di Filosofia e Psicologia* 7: 256–263.

- Searle, J. (1984) *Minds, Brains and Science*, Cambridge, MA: Harvard University Press.
- Searle, J. (1987) "Indeterminacy, Empiricism, and the First Person," *Journal of Philosophy* 84: 123–146.
- Searle, J. (1992) *The Rediscovery of the Mind*, Cambridge, MA: MIT Press.
- Siewert, C. (1998) *The Significance of Consciousness*, Princeton, NJ: Princeton University Press.
- Siewert, C. (2011) "Phenomenal Thought," in Bayne and Montague (2011).
- Sterelny, K. (1990) *The Representational Theory of Mind*, Oxford: Blackwell.
- Strawson, G. (1986) *Freedom and Belief*, Oxford: Oxford University Press.
- Strawson, G. (1994) *Mental Reality*, Cambridge, MA: MIT Press.
- Strawson, G. (2008) "Real Intentionality 3: Why Intentionality Entails Consciousness," in *Real Materialism and Other Essays*, Oxford: Oxford University Press.
- Strawson, G. (2011) "Cognitive Phenomenology: Real Life," in Bayne and Montague (2011).
- Thompson, B. (2008) "Representationalism and the Argument from Hallucination," *Pacific Philosophical Quarterly* 89: 384–412.
- Tye, M. (2000) *Consciousness, Color and Content*, Cambridge: MIT Press.
- Tye, M. and Wright, B. (2011) "Is There a Phenomenology of Thought?" in Bayne and Montague (2011).

Related Topics

Representational Theories of Consciousness
Consciousness and Conceptualism
Dualism
Materialism